

Predicting wins in high ELO league of legends

Antonio Kristófer Salvador
Axel Magnússon



Figure 1: League of Legends


1 The dataset

League of Legends is a 2009 multiplayer online battle arena video game developed and published by Riot Games where two teams (red and blue) of 5 players fight to destroy their opponents base. The dataset contains the stats at the ten minute mark for close to ten thousand games from high ELO (ranks from diamond 1 to master) - approximately the top 0.3 percent of players. Our goal in this project is to use these statistics to predict the end result of the game. There are 19 features per team, a unique gameld and the target value blueWins. The dataset comes from [here](#) and there are some more details along with a short glossary on some league of legends terms the reader may find useful.

```
RangeIndex: 9879 entries, 0 to 9878
Data columns (total 40 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   gameId              9879 non-null   int64
1   blueWins            9879 non-null   int64
2   blueWardsPlaced     9879 non-null   int64
3   blueWardsDestroyed  9879 non-null   int64
4   blueFirstBlood      9879 non-null   int64
5   blueKills           9879 non-null   int64
6   blueDeaths          9879 non-null   int64
7   blueAssists         9879 non-null   int64
8   blueEliteMonsters   9879 non-null   int64
9   blueDragons         9879 non-null   int64
10  blueHeralds         9879 non-null   int64
11  blueTowersDestroyed  9879 non-null   int64
12  blueTotalGold       9879 non-null   int64
13  blueAvgLevel        9879 non-null   float64
14  blueTotalExperience  9879 non-null   int64
15  blueTotalMinionsKilled  9879 non-null   int64
16  blueTotalJungleMinionsKilled  9879 non-null   int64
17  blueGoldDiff        9879 non-null   int64
18  blueExperienceDiff   9879 non-null   int64
19  blueCSPerMin        9879 non-null   float64
20  blueGoldPerMin      9879 non-null   float64
21  redWardsPlaced      9879 non-null   int64
22  redWardsDestroyed   9879 non-null   int64
23  redFirstBlood       9879 non-null   int64
24  redKills            9879 non-null   int64
25  redDeaths           9879 non-null   int64
26  redAssists          9879 non-null   int64
27  redEliteMonsters    9879 non-null   int64
28  redDragons          9879 non-null   int64
29  redHeralds          9879 non-null   int64
30  redTowersDestroyed  9879 non-null   int64
31  redTotalGold        9879 non-null   int64
32  redAvgLevel         9879 non-null   float64
33  redTotalExperience   9879 non-null   int64
34  redTotalMinionsKilled  9879 non-null   int64
35  redTotalJungleMinionsKilled  9879 non-null   int64
36  redGoldDiff         9879 non-null   int64
37  redExperienceDiff    9879 non-null   int64
38  redCSPerMin         9879 non-null   float64
39  redGoldPerMin       9879 non-null   float64
dtypes: float64(6), int64(34)
memory usage: 3.0 MB
```

2 Preprocessing

To begin with preprocessing we first looked to see if any values were null or otherwise invalid, we then took note of the correlation between each feature and the target feature.




A few features were obviously redundant, having a correlation of one with other features, those features being blueTotalGold and blueGoldPerMin because all games are exactly 10 minutes in and therefore gold/min is simply totalgold/10 same for totalminion-skilled and cs/minute (cs stands for creep score - the creeps being minions). So we removed those time based attributes, along with the gameld as it had no correlations.

The feature "wards placed", had us confused, the values being far too high, we suspect this column is incorrectly labelled and refers to vision score, which measures effectiveness of wards placed, and destroyed (+1 for each minute of ward lifetime gained and +1 for each minute of ward lifetime denied), however in the end we decided to keep it. both sides have a gold and exp diff column, which are inverses of each other and so we removed the diff attributes from the red side.

The values with the highest correlations with our target value were: Golddiff, Expdiff, Totalgold, Totalexperience, Avglevel, Kills and Deaths.

We also removed the attributes redKills and redDeaths because the red teams kills are represented in the blue teams deaths and vice versa.

We tried removing the outliers by filtering out all values outside of three standard deviations, but later decided against it since the outliers are based on two cases: One team is dominating and should therefore be likely to win or one team or some members are deliberately losing, this can be seen in wardsplaced since the only way to have scores like the 250+ ones we observed is to sit in the teams base and use up all of your gold on wards.



We also tried making delta variables for all variables that had both blue and red counterparts but in the end we decided against it since we felt we were losing some information along the way, for example a team leading 14-10 in kills is not as significant as a team leading 4-0 but both would give us a delta value of 4.

3 Models

After preprocessing we decided on which models to try. These are the models we chose: DecisionTreeClassifier, SelectKBest, MLPClassifier, SVC, RandomForestClassifier, LogisticRegression, KNeighborsClassifier

DecisionTreeClassifier: Because we used this in a lab assignment and were somewhat familiar with it. We also did a second DecisionTreeClassifier but using SelectKBest for comparison.

We tried the MLPClassifier because we wanted to try working with neural networks.

Logistic Regression: This algorithm is used for binary classification problems which is exactly what we have so we thought this could work well.

SVC: Works well with 2 classes like logistic regression, is easy to

use and unlike decision trees they tend to not overfit data.


All of our models were trained using GridSearchCV in conjunction with sklearn pipelines, this allowed us to easily test many permutations of different hyperparameters for each model efficiently, while providing cross validation simultaneously.

4 Results


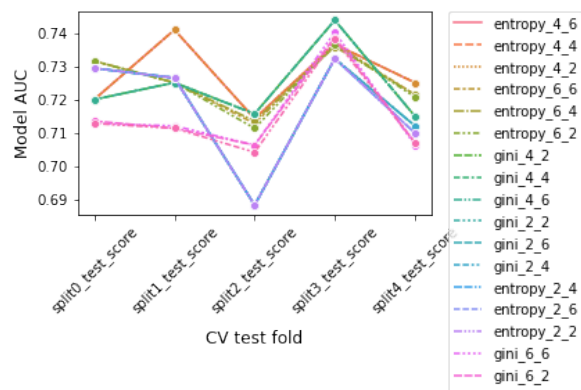
4.1 introduction

We include a confusion matrix and a classification report for the best model of each type.

4.2 Decision tree



	precision	recall	f1-score	support
blue loses	0.72	0.73	0.73	1506
blue wins	0.72	0.71	0.72	1458
accuracy			0.72	2964
macro avg	0.72	0.72	0.72	2964
weighted avg	0.72	0.72	0.72	2964







Figure 2: min samples split ranging from 2 - 1000


4.3 SelectKBest



	precision	recall	f1-score	support
blue loses	0.72	0.73	0.73	1506
blue wins	0.72	0.71	0.72	1458
<hr/>				
accuracy			0.72	2964
macro avg	0.72	0.72	0.72	2964
weighted avg	0.72	0.72	0.72	2964

The K value had very little effect, results were largely similar to the regular decision tree.


4.4 MLPClassifier




	precision	recall	f1-score	support
blue loses	0.73	0.73	0.73	1506
blue wins	0.72	0.72	0.72	1458
<hr/>				
accuracy			0.72	2964
macro avg	0.72	0.72	0.72	2964
weighted avg	0.72	0.72	0.72	2964

The hidden layer sizes were rather confusing, we believe we hit somewhat of a sweetspot at around (70,) however this may not be the best value, we see a peak at (30,) which corresponds with our number of attributes, the greater values may simply be approximating more complex networks in a single


layer, we also tested multi layer setups, but failed to find one that performed better than their singly layered counterparts.




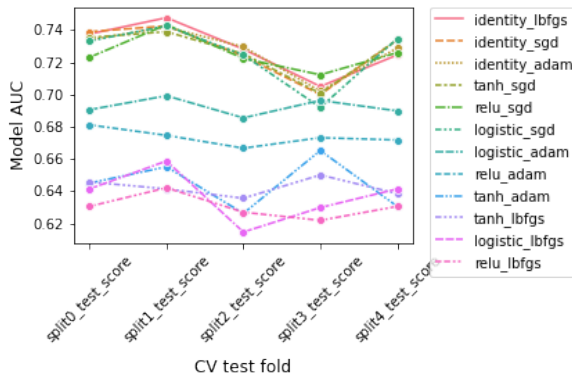
Identity was not as consistently the best performing activation function as this graph depicts, the rankings seemed to vary based on the hidden layer sizes. However relu and logistic were consistently the worst.



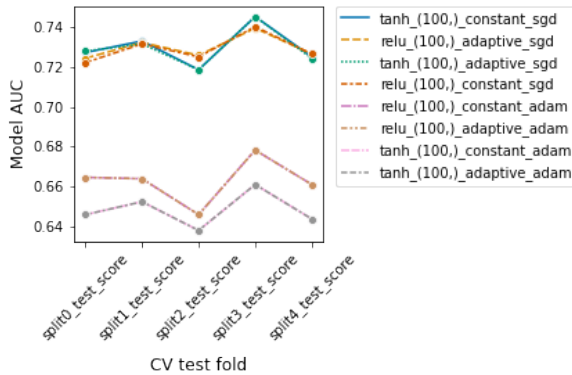
As for learning rates invscaling was consistently the worst performing, but the others were mixed.



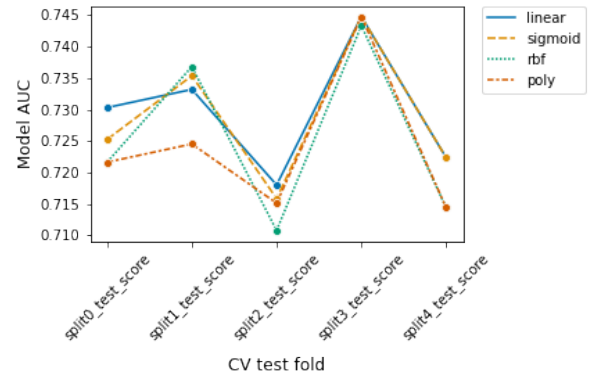
Solver: sgd was best most of the time, except when using the identity activation.




This next graph shows how the performance of different activation functions relied more on the other values as the ordering is inconsistent. It also shows how sgd was consistently better than adam.



The kernels were somewhat similar in performance, but linear and sigmoid were on top for most tests. This leads us to believe our dataset is at least somewhat linearly separable.




4.5 SVC



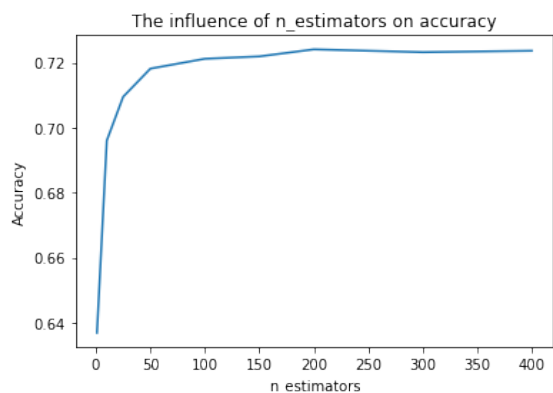
	precision	recall	f1-score	support
blue loses	0.72	0.73	0.73	1506
blue wins	0.72	0.71	0.72	1458
<hr/>				
accuracy			0.72	2964
macro avg	0.72	0.72	0.72	2964
weighted avg	0.72	0.72	0.72	2964

Smaller C values consistently outperformed larger ones, smaller values here being values lower than one. We believe this is because the dataset is rather complex but still well separable if we give the model enough freedom. We also tested extremely small values and this trend continued, the best performing tested value being 0.001, which makes sense as the data is quite noisy, since many games are either not decided yet at ten minutes or a team may blow their lead and lose.

4.6 Random Forest Classifier




	precision	recall	f1-score	support
blue loses	0.73	0.73	0.73	1506
blue wins	0.72	0.72	0.72	1458
<hr/>				
accuracy			0.73	2964
macro avg	0.73	0.73	0.73	2964
weighted avg	0.73	0.73	0.73	2964




This graph shows the influence of the number of trees in the forest. We tested values ranging from 1-400 and as we can see the graph stabilizes and we start to have diminishing returns at approximately $n=100$.

The influence of randomforestclassifier_min_samples_split on accuracy



4.7 Logistic Regression




	precision	recall	f1-score	support
blue loses	0.73	0.73	0.73	1506
blue wins	0.72	0.72	0.72	1458
<hr/>				
accuracy			0.72	2964
macro avg	0.72	0.72	0.72	2964
weighted avg	0.72	0.72	0.72	2964


For logistic regression we tried changing the solver algorithm but this ended up having negligible effect on accuracy so we stuck with the default "lbfgs".

4.8 KNeighborsClassifier

This graph is almost identical to the one in the RFC. We see a big jump in accuracy gain from 1-100 but then it stabilizes.



	precision	recall	f1-score	support
blue loses	0.72	0.72	0.72	1506
blue wins	0.71	0.71	0.71	1458
<hr/>				
accuracy			0.72	2964
macro avg	0.72	0.72	0.72	2964
weighted avg	0.72	0.72	0.72	2964



5 Conclusions

All of our models ended up having roughly the same accuracy, around 0.71 - 0.73. We believe this is close to the limit we can reach with the training data we are working with. The model which was selected was not always of the same type, we believe this is a result of all of the models being so close in performance and the best one decides on random factors.

6 Future work

We might have achieved a better result by focusing more of our efforts on the feature engineering side, we would spend more time on this next time. We should also set the random state attribute on all models to have more concrete performance comparisons.